


# LSTM and GRU

Debapriyo Majumdar

[debapriyo@isical.ac.in](mailto:debapriyo@isical.ac.in)

# Shortcomings of vanilla RNN

2



- The network cannot remember “enough” if the sequence is too long
- *Example:* Samuel **spent his childhood in Spain**. Then he lived in France, Germany and England. However, **he can still speak Spanish**.

# Shortcomings of vanilla RNN: vanishing gradient


3

Activation to the next cell  $\rightarrow$

$$a^{<t>} = g(W_a a^{<t-1>} + W_x x^{<t>} + b)$$


$\leftarrow$  Backpropagation from the next cell

$$\frac{\partial L}{\partial a^{<t-1>}} = g' \cdot W_a^T \frac{\partial L}{\partial a^{<t>}}$$



## Vanishing gradient problem

- Long chain of *multiplication* with  $W_a^T$  and  $g'$  (derivative of the activation)
- Both tanh and sigmoid have small derivatives
- ReLU is not used in RNN because the values can explode



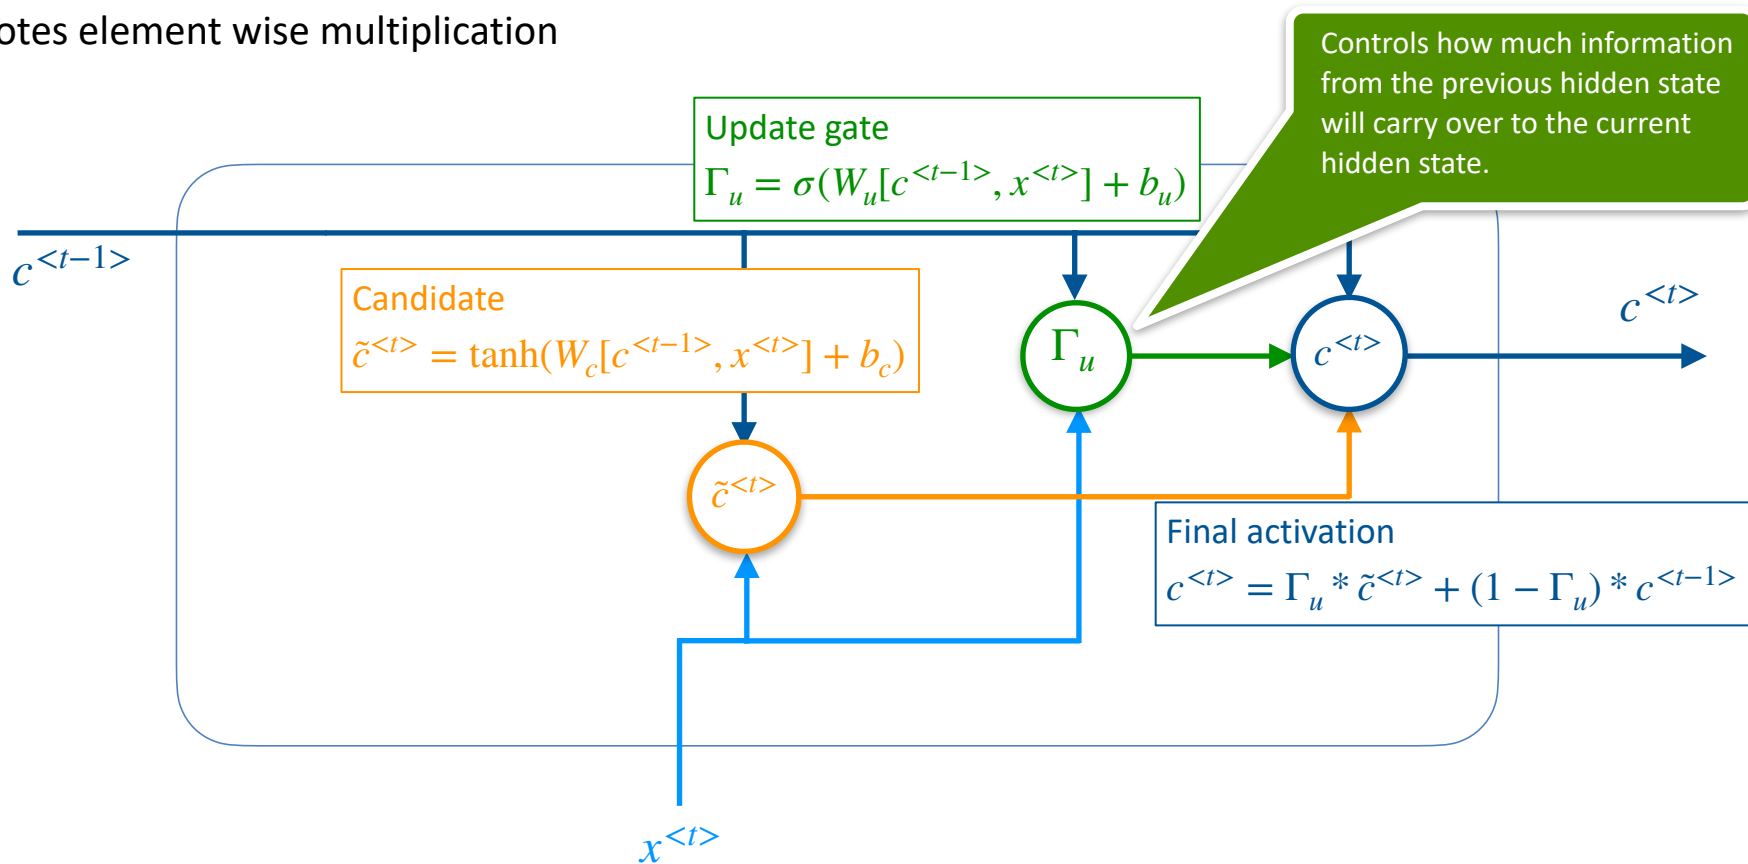
# Gated Recurrent Unit (GRU): Simplified

4

Notation: memory cell  $c^{<t>}$

For GRU,  $c^{<t>}$  is same as activation  $a^{<t>}$  as before

\* denotes element wise multiplication



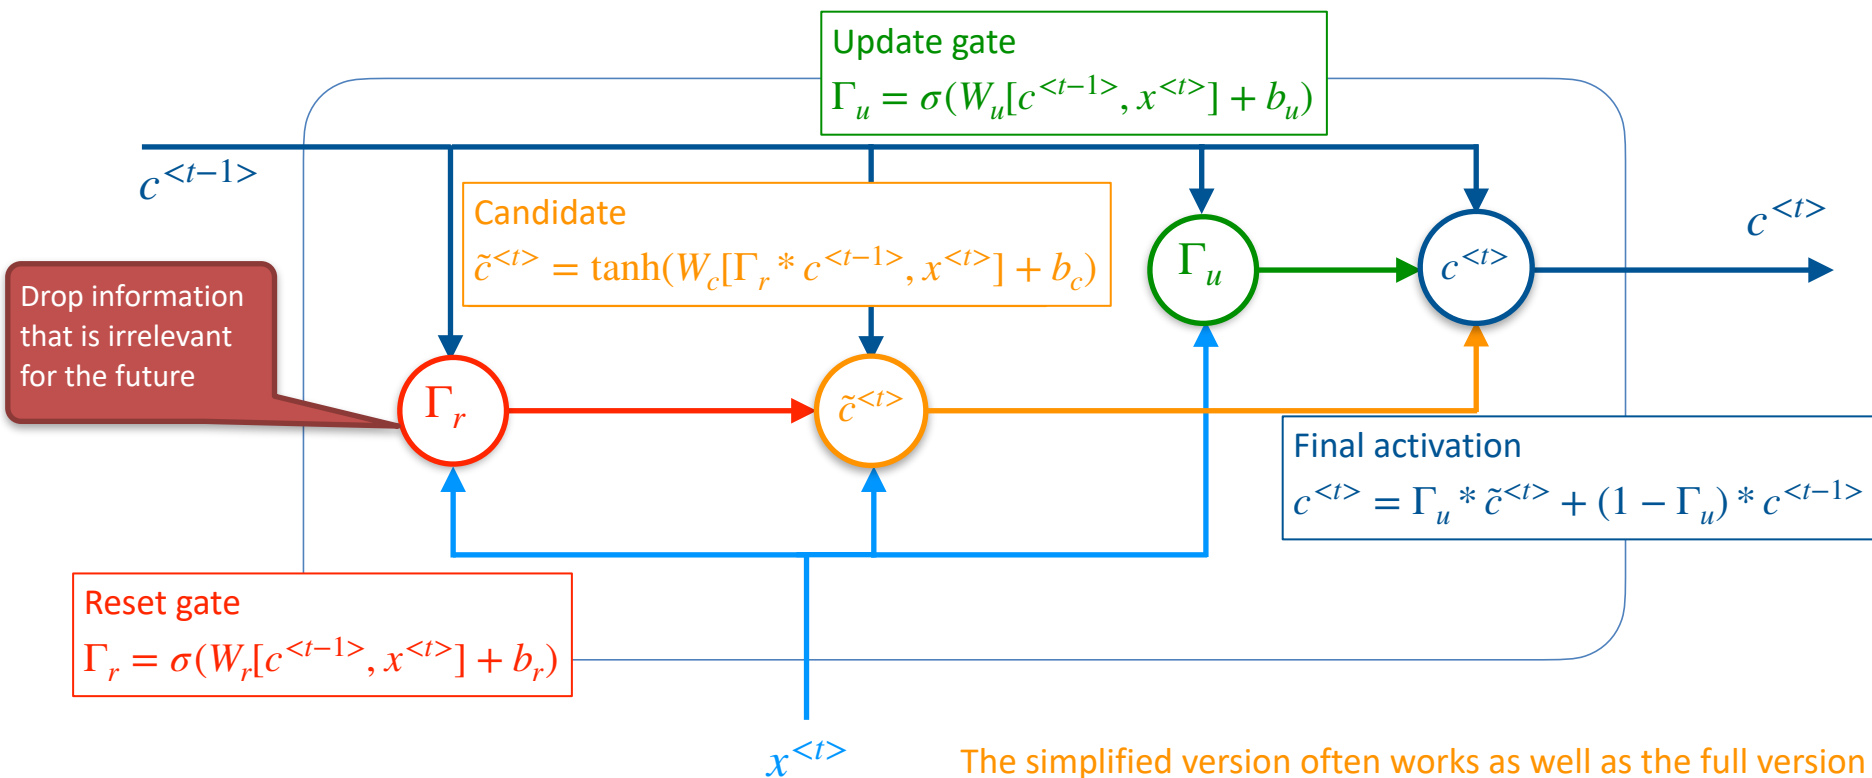
# Gated Recurrent Unit (GRU)

5

Notation: memory cell  $c^{<t>}$

For GRU,  $c^{<t>}$  is same as activation  $a^{<t>}$  as before

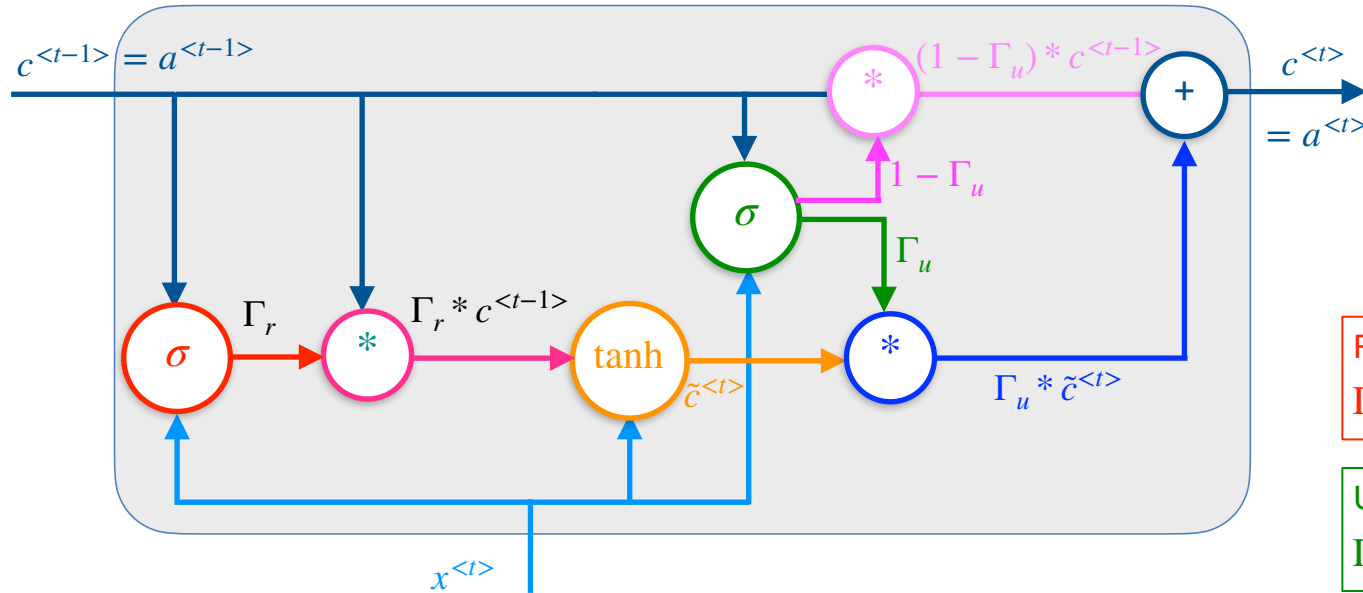
\* denotes element wise multiplication



# Gated Recurrent Unit (GRU)

6

Notation: memory cell  $c^{<t>}$  (for GRU,  $c^{<t>} = \text{activation } a^{<t>}$ ),  $*$  denotes element wise multiplication



Reset gate

$$\Gamma_r = \sigma(W_r[c^{<t-1>}, x^{<t>}] + b_r)$$

Update gate

$$\Gamma_u = \sigma(W_u[c^{<t-1>}, x^{<t>}] + b_u)$$

Candidate

$$\tilde{c}^{<t>} = \tanh(W_c[\Gamma_r * c^{<t-1>}, x^{<t>}] + b_c)$$

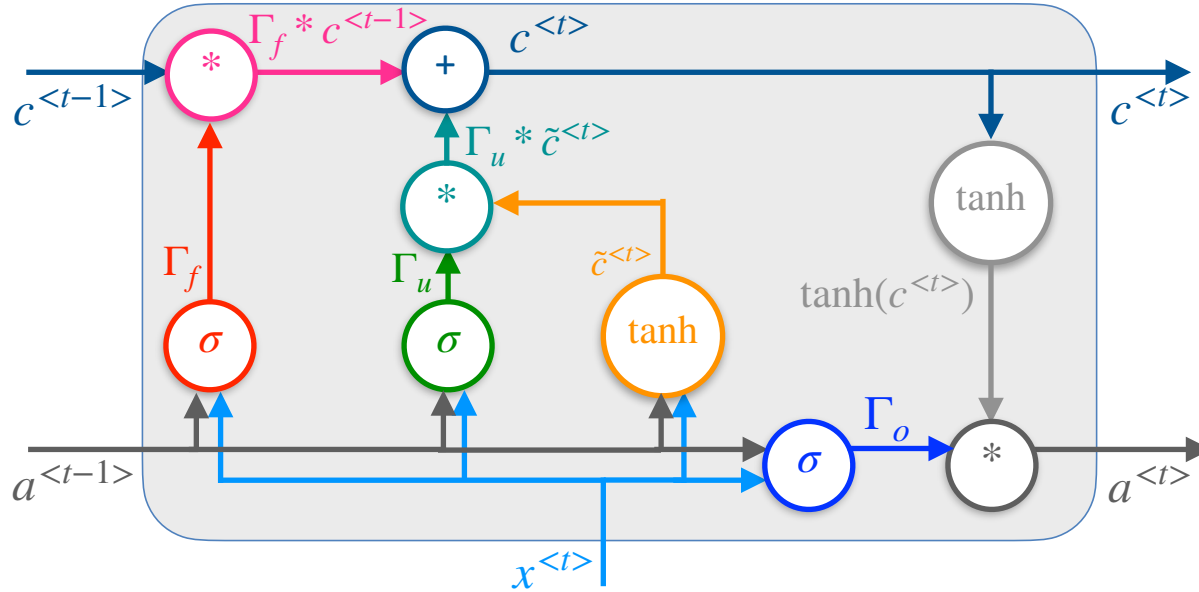
Final activation to the next cell

$$c^{<t>} = \Gamma_u * \tilde{c}^{<t>} + (1 - \Gamma_u) * c^{<t-1>}$$

# Long Short-Term Memory (LSTM)

7

Notation: memory cell  $c^{<t>}$ , activation  $a^{<t>}$ , \* denotes element wise multiplication



Forget gate

$$\Gamma_f = \sigma(W_f[a^{<t-1>}, x^{<t>}] + b_f)$$

Update gate

$$\Gamma_u = \sigma(W_u[a^{<t-1>}, x^{<t>}] + b_u)$$

Candidate activation

$$\tilde{c}^{<t>} = \tanh(W_c[a^{<t-1>}, x^{<t>}] + b_c)$$

Memory to the next cell

$$c^{<t>} = \Gamma_u * \tilde{c}^{<t>} + \Gamma_f * c^{<t-1>}$$

Output gate

$$\Gamma_o = \sigma(W_o[a^{<t-1>}, x^{<t>}] + b_o)$$

Activation to the next cell

$$a^{<t>} = \Gamma_o * \tanh(c^{<t>})$$

- Memory from the previous cell pass through the gate  $\Gamma_f$

# LSTM vs GRU

8

LSTM	GRU
More widely used	Recently gained popularity
More powerful	Simpler, but works almost as well as LSTM
Computationally expensive	Computationally cheaper, so deeper networks can be trained



- Andrew Ng's lectures on *Sequence Models*: [www.coursera.org/learn/nlp-sequence-models](http://www.coursera.org/learn/nlp-sequence-models)
- Chris Manning, Abigail See and other TAs. *Natural Language Processing with Deep Learning*. Stanford University Course (CS224n), Winter 2019. [web.stanford.edu/class/archive/cs/cs224n/cs224n.1194/](http://web.stanford.edu/class/archive/cs/cs224n/cs224n.1194/)
- Goodfellow, Ian, Yoshua Bengio, and Aaron Courville. *Deep learning*. MIT press, 2016. [www.deeplearningbook.org](http://www.deeplearningbook.org)
- Sepp Hochreiter and Jürgen Schmidhuber. 1997. Long Short-Term Memory. *Neural Comput.* 9, 8 (November 15, 1997), 1735–1780. DOI:<https://doi.org/10.1162/neco.1997.9.8.1735>
- Nir Arbel. *How LSTM networks solve the problem of vanishing gradients*. 2018. [medium.com/datadriveninvestor/how-do-lstm-networks-solve-the-problem-of-vanishing-gradients-a6784971a577](https://medium.com/datadriveninvestor/how-do-lstm-networks-solve-the-problem-of-vanishing-gradients-a6784971a577)
- Cho, Kyunghyun, Bart Van Merriënboer, Caglar Gulcehre, Dzmitry Bahdanau, Fethi Bougares, Holger Schwenk, and Yoshua Bengio. "[Learning phrase representations using RNN encoder-decoder for statistical machine translation](https://arxiv.org/abs/1406.1078)." *arXiv preprint arXiv:1406.1078* (2014).